Career Goals & Future research

MPhil in Data Intensive Science
Samvardhan Singh

My long-term career goal is to become a leading AI Researcher in the Physical Sciences, focusing on the development of verifiable, high-precision discovery frameworks for exascale data. Having architected production-scale MLOps at NielsenIQ and led institutional tech initiatives at the IIE Center, I aim to lead international research labs that bridge the gap between industrial reliability and scientific discovery.
I wish to propose a hybrid research project that utilizes Wavelet-based Signal Decomposition and Quantum Feature Mapping to detect rare astronomical transients, such as Fast Radio Bursts (FRBs) or gravitational wave anomalies within high-noise environments, if given a chance, it is my aim to work on such research projects during my masters:

Quantum-Enhanced Transient Discovery (QETD): A Hybrid Architecture for High-Noise Astronomical Anomaly Detection

Executive Summary

The impending operational commencement of the Square Kilometre Array (SKA) and the High-Luminosity Large Hadron Collider (HL-LHC) marks the transition of physical sciences into the Exascale Era. This shift is characterized not merely by the volume of data projected in the exabytes per year but by the complexity of the signals buried within. We stand on the precipice of discovering the universe's most elusive phenomena: Fast Radio Bursts (FRBs) that probe the intergalactic medium, gravitational wave anomalies revealing exotic compact object mergers, and potentially, technosignatures of extraterrestrial origin. However, the sensitivity required to detect these faint, transient events simultaneously exposes our detectors to an overwhelming deluge of anthropogenic and terrestrial noise.
Current signal processing paradigms, predicated on classical matched filtering and Fourier-domain thresholding, are approaching their asymptotic limits. They are computationally expensive, require rigid templates that blind us to "unknown unknowns," and struggle to distinguish non-Gaussian astrophysical transients from non-stationary Radio Frequency Interference (RFI). Deep learning offers a reprieve but faces its own "plateau of diminishing returns" regarding training data requirements and feature extraction in low-signal-to-noise (SNR) regimes.
This report articulates a comprehensive research proposal for Quantum-Enhanced Transient Discovery (QETD), a hybrid computing architecture designed to transcend these limitations. QETD synergizes Wavelet-based Signal Decomposition, specifically the Wavelet Scattering

Transform with Quantum Feature Mapping to create a detection pipeline of unprecedented sensitivity and specificity. The core hypothesis is that while classical algorithms excel in Euclidean feature spaces, the complex, entangled temporal correlations of transient signals are better represented in the high-dimensional Hilbert spaces accessible via quantum entanglement.
By preprocessing terabit-per-second data streams using scale-invariant wavelet transforms to reduce dimensionality while preserving signal morphology, and subsequently mapping these features into a quantum kernel, the system achieves a separation of signal and noise mathematically impossible in lower-dimensional classical spaces. This document provides an exhaustive technical analysis of the problem, the proposed hybrid solution, the mathematical foundations of wavelet theory and quantum kernels, and the requisite MLOps infrastructure to deploy such a system at the scale of the SKA Regional Centres. It outlines a phased development roadmap from simulation to exascale deployment, arguing that the integration of quantum intelligence into astrophysical workflows is a necessity for the next frontier of discovery.

1. Introduction: The Exascale Horizon and the Transient Sky

1.1 The Paradigm Shift to Time-Domain Astronomy

For much of its history, astronomy was a static science, concerned with mapping the fixed positions of stars and galaxies. The 21st century has ushered in the era of Time-Domain Astronomy, where the universe is understood as a dynamic, bursting, and evolving entity. Events occur on timescales ranging from fractions of a second to years. Among these, the most enigmatic are millisecond-duration radio transients.

The discovery of Fast Radio Bursts (FRBs), intense pulses of radio emission lasting only milliseconds yet releasing as much energy as the Sun does in days, has revolutionized high-energy astrophysics. Since the first detection in 2007, FRBs have become a premier probe of the intergalactic medium (IGM), offering unique insights into the "missing baryons" of the universe and the fundamental physics of extreme magnetic fields.[1] Similarly, the detection of Gravitational Waves (GW) by LIGO and Virgo has opened a non-electromagnetic window onto the universe, allowing us to hear the chirps of merging black holes and neutron stars.[3]

However, the transient sky is fleeting. Unlike a galaxy that can be re-observed indefinitely, a transient event is lost forever if not captured in real-time. This imposes a strict "deadline for discovery." The detection system must identify, classify, and trigger follow-up observations within seconds of the event's arrival at the detector.

1.2 The Data Deluge: SKA and LHC Challenges

The scale of data generation in modern physics is staggering. The Square Kilometre Array (SKA), currently under construction in Australia and South Africa, will be the largest radio telescope ever built.

- **Data Rates:** The SKA-Mid telescope alone is projected to produce raw data rates from the correlator exceeding 8.9 Tb/s.[4]
- **Volume:** High-priority science objectives are expected to produce 300 PB of data per telescope per year during full operations.[5]
- **Processing:** The Science Data Processor (SDP) must ingest this stream, perform calibration, imaging, and transient search in near real-time.

A parallel challenge exists at the Large Hadron Collider (LHC). The detectors there operate as "cameras" taking 40 million pictures per second. The trigger system must discard 99.998% of these events immediately, keeping only the most "interesting" ones for analysis. As the LHC moves to High-Luminosity operation, the pile-up of background events (noise) increases effectively, necessitating smarter, faster trigger algorithms.[3]

In both domains, the challenge is **Data Intensive Science**. We are searching for rare anomalies (a few events per day or week) in a stream of zeros and ones flowing at terabits per second.

1.3 The Failure of Classical Approaches

The traditional workhorse of transient detection is **Matched Filtering**. This involves convolving the incoming data stream with a bank of pre-computed templates (theoretical waveforms of FRBs or GWs).[3]

- **Computational Cost:** The number of templates required to cover the parameter space (dispersion measure, pulse width, scattering, chirp mass) is massive. For gravitational waves, template banks can contain hundreds of thousands of waveforms.
- **Blindness:** Matched filtering is optimal only for *known* signals embedded in Gaussian noise. It is blind to "unknown unknowns"; signals that do not match the templates.
- **RFI Susceptibility:** Terrestrial Radio Frequency Interference (RFI) is non-Gaussian and often structurally similar to astrophysical signals. A classical filter bank produces millions of false positives when faced with complex RFI environments.[2]

Recent advances in Deep Learning (DL), such as the DRAFTS pipeline (Deep learning-based RAdio Fast Transient Search), have utilized Convolutional Neural Networks (CNNs) to improve detection rates.[7] While effective, these models require massive labeled training datasets, which do not exist for undiscovered phenomena. Furthermore, as RFI environments become more complex (with satellite constellations like Starlink), the complexity of the neural networks must increase, leading to latency issues in real-time inference.

1.4 The QETD Proposition

This research proposes **Quantum-Enhanced Transient Discovery (QETD)** as a solution to the sensitivity and specificity bottlenecks of classical systems. The core premise is that quantum computers, operating in Hilbert space, can access feature maps that are computationally intractable for classical computers. By encoding astronomical data into quantum states, we can exploit entanglement to capture high-order correlations in the time-series data, enabling the separation of subtle astrophysical signals from complex, structured background noise.[9]

The QETD architecture is hybrid:

1. **Classical Preprocessing:** Utilizing **Wavelet Scattering Transforms** to perform robust, deformation-invariant feature extraction on the raw high-rate data.[11]
2. **Quantum Classification:** Utilizing **Quantum Kernels** and **Variational Quantum Circuits** to classify these features with high precision and low false-alarm rates, even with limited training data (Quantum Advantage in sample complexity).[13]

This report details the theoretical, mathematical, and engineering path to realizing this architecture.

2. The Scientific Context: Anatomy of Signals and Noise

To design a superior detection system, one must possess a nuanced understanding of the signal's morphology and the statistical properties of the noise that obscures it.

2.1 Astrophysical Transients: The Chirp and the Burst

The two primary targets of this research, Fast Radio Bursts (FRBs) and Gravitational Waves (GWs) share a fundamental characteristic: they are **dispersed or chirped transients**.

2.1.1 Fast Radio Bursts (FRBs)

FRBs are millisecond-duration radio pulses. Their defining feature is **Dispersion**. As the radio wave travels through the ionized plasma of the Intergalactic Medium (IGM), interaction with free electrons causes a frequency-dependent time delay. Lower frequencies arrive later than higher frequencies. The time delay Δt follows the cold plasma dispersion relation:

$$\Delta t \propto \mathrm{DM} \times (\nu_{\mathrm{low}}^{-2} - \nu_{\mathrm{high}}^{-2})$$

where DM (Dispersion Measure) is the integrated column density of free electrons along the line of sight.[2]

- **De-dispersion Challenge:** To detect an FRB, the detector must correct for this delay. Since the DM is unknown beforehand, the system must trial thousands of potential DM values, summing the energy across frequency channels for each trial. This is the most computationally intensive step in current pipelines.[4]
- **Scattering:** In addition to dispersion, multi-path propagation through turbulent plasma causes "scattering," broadening the pulse exponentially (scattering tail). This asymmetry is a key morphological marker distinguishing real FRBs from symmetric RFI bursts.

2.1.2 Gravitational Wave Anomalies

Gravitational waves from compact binary coalescences (black holes or neutron stars) manifest as a "chirp" signal; a waveform that sweeps up in both frequency and amplitude as the objects spiral inward.

$$h(t) \propto A(t)\cos(\Phi(t))$$

While the matched filter approach works well for these standard "inspirals," it struggles with:

- **High-Mass Mergers:** Short-duration signals from intermediate-mass black holes that have very few cycles in the detector's sensitivity band.[3]
- **Eccentric Binaries:** Systems with elliptical orbits produce complex waveforms that standard templates miss.
- **Glitches:** The detectors suffer from instrumental noise transients called "glitches" (e.g., "blips," "scattered light") that can mimic short GW signals, creating a high background of false alarms.[3]

2.2 The Noise Landscape: A Structural Problem

The "noise" in these experiments is not merely thermal hiss; it is structural and adversarial.

2.2.1 Thermal Noise (Gaussian)

This is the irreducible background caused by the random thermal motion of electrons in the receiver system. It is statistically Gaussian and "white" (flat power spectrum) over the bandwidth.[15] Classical signal processing theory (Shannon-Nyquist, Matched Filtering) is built to handle this type of noise optimally.

2.2.2 Radio Frequency Interference (RFI) (Non-Gaussian)

The primary enemy of the radio astronomer is RFI. It originates from human activity: mobile phones, digital TV, satellites, electric fences, and even the telescope's own electronics.

- **Narrowband RFI:** Persistent signals at specific frequencies (e.g., GPS, GSM). These appear as vertical lines in a Time-Frequency spectrogram.[2]
- **Broadband RFI:** Impulsive bursts from lightning or switching power supplies. These appear as horizontal lines or blobs in the spectrogram.

- **Wandering RFI:** Signals that drift in frequency, potentially mimicking the ν^{-2} dispersion sweep of an FRB.

The statistical distribution of RFI is **Heavy-Tailed** (non-Gaussian). A 5-sigma event in Gaussian noise is a rare discovery; a 5-sigma event in RFI-dominated data happens every few seconds. This breakdown of Gaussian statistics is why simple thresholding fails and why advanced AI is needed.[2]

2.3 Cross-Disciplinary Insight: The Neuroscience Connection

A critical insight driving the QETD proposal is the parallel between radio astronomy and neuroscience.

- **EEG Analysis:** Electroencephalograms (EEG) measure millisecond-scale electrical spikes from the brain (neural oscillations) embedded in a sea of biological and environmental noise (muscle movement, 60Hz power line hum).[17]
- **Parallels:** Both fields deal with identifying non-stationary, transient oscillatory events (epileptic spikes vs. FRBs) in high-noise time-series data.
- **Lesson Learned:** Neuroscience moved away from pure Fourier analysis (which loses time information) to **Wavelet Analysis** to successfully identify these transients.[18] Wavelets allow researchers to isolate "spindles" and "spikes" that are invisible in the averaged power spectrum.

QETD applies this biological signal processing philosophy to the cosmos: treating the radio sky not as a static image, but as a dynamic "neural" signal that requires time-frequency decomposition to decode.[20]

3. The Computational Bottleneck: Why Classical AI is Insufficient

Before introducing the quantum solution, it is necessary to articulate why current "Classical AI" (Deep Learning) is nearing a performance plateau.

3.1 The Limits of Convolutional Neural Networks (CNNs)

Pipelines like DRAFTS [7] and others using architectures like ResNet or CenterNet [8] treat the Time-Frequency spectrogram as an image.

- **Inductive Bias:** CNNs possess translation invariance, which is good. However, they struggle with the specific *deformation* invariance required for FRBs. A highly dispersed FRB looks stretched compared to a low-DM FRB. Standard CNNs must learn this relationship by seeing thousands of examples of every possible stretch.

- **Data Hunger:** Training a robust CNN requires massive datasets. For FRBs, we have only ~800 real confirmed sources.[2] To train models, we rely on simulations (injected fake signals). If the simulations do not perfectly match reality (e.g., complex scattering, scintillation), the model fails to generalize to real data (sim-to-real gap).
- **The "Black Box" Risk:** In science, interpretability is key. A CNN that outputs "Probability 0.99" without explanation is dangerous. It might be triggering on a subtle RFI feature that correlates with the signal in the training set (the "Clever Hans" effect).

3.2 The Feature Space Problem

Classical Support Vector Machines (SVMs) and Kernel methods rely on mapping data to a higher-dimensional feature space to find a linear separation boundary.

$$K(x, x') = \langle \phi(x), \phi(x') \rangle$$

Common kernels like the Radial Basis Function (RBF) or Polynomial kernel are computationally efficient but limited in expressivity. They effectively map data to a space that is still fundamentally constrained by classical correlations.

In high-noise environments, the "signal manifold" and the "noise manifold" might be inextricably enshrined in classical Euclidean space. No amount of classical non-linear transformation can separate them if they overlap in all accessible dimensions.

3.3 The Latency Budget

The SKA is a streaming instrument. The data flow is continuous.

- **Inference Time:** A deep ResNet-50 might take 50-100ms to process a spectrogram on a GPU. If the data rate requires processing every 10ms, the system falls behind.
- **Power Consumption:** Running massive GPU clusters for inference at the edge (in the remote Australian desert) is power-prohibitive. We need algorithms that are computationally dense but energy efficient (or that can offload the complexity to a quantum state).

4. Proposed Solution Part I: Wavelet-Based Signal Decomposition

The first pillar of the QETD architecture is a mathematically rigorous preprocessing stage designed to prepare the data for the quantum classifier. We reject the standard FFT in favor of the **Wavelet Scattering Transform (WST)**.

4.1 Beyond the Fourier Transform

The Fast Fourier Transform (FFT) assumes signal stationarity, that the frequency content does not change over time. Transients are, by definition, non-stationary. A short burst, when FFT'd, spreads its energy across the entire phase space, disappearing into the noise floor.[20] The Short-Time Fourier Transform (STFT) attempts to fix this by windowing, but it is bound by the Heisenberg-Gabor uncertainty limit: one cannot simultaneously have perfect time and frequency resolution. Fixed windows (like in STFT) force a compromise that is often suboptimal for both short bursts (FRBs) and long chirps (GWs).[18]

4.2 The Wavelet Scattering Transform (WST)

Developed by Stéphane Mallat, the WST is often described as a "Deep Neural Network with fixed, mathematical weights".[12] It provides a representation that is stable to deformations and translation-invariant. Unlike a standard Wavelet Transform (which is linear and variant to shifts), the WST introduces non-linearity via the **Modulus Operator**.

The Cascade Mechanism:

1. **Zeroth Order:** The average of the signal (DC component).

$$S_0 x = x * \phi$$

2. **First Order:** Wavelet convolution followed by modulus and averaging. This captures the energy in different frequency bands (scalogram).

$$S_1 x(t, \lambda) = |x * \psi_\lambda| * \phi(t)$$

3. **Second Order:** This is the crucial step for transients. It takes the *modulus* of the first layer and convolves it *again* with a second bank of wavelets.

$$S_2 x(t, \lambda_1, \lambda_2) = ||x * \psi_{\lambda_1}| * \psi_{\lambda_2}| * \phi(t)$$

Why WST for Astronomy?

- **Interference Cancellation:** The WST separates processes based on their scale and modulation. Gaussian thermal noise tends to scatter energy uniformly. Coherent RFI and astrophysical bursts scatter energy into specific "paths" (λ_1, λ_2) in the scattering network.[22]
- **Cosmological Sensitivity:** Research in the 21cm forest has shown that WST captures non-Gaussian information about the Intergalactic Medium that the power spectrum completely misses.[11]
- **Data Compression:** The WST reduces a high-dimensional time-series (e.g., 4096 voltage samples) into a compact coefficient vector (e.g., 64 coefficients) that retains the *morphological essence* of the signal. This compression is vital for enabling Quantum Machine Learning, which currently has limited input bandwidth.[24]

In QETD, the WST acts as the "Classical Compressor," discarding the bulk of the Gaussian noise and presenting the Quantum Computer with a distilled feature vector rich in structural information.

5. Proposed Solution Part II: Quantum Feature Mapping and Classification

The second pillar of QETD is the utilization of Quantum Machine Learning (QML) to classify the WST-derived features. This leverages the "Hilbert Space Advantage."

5.1 The Quantum Kernel Trick

Classical kernel methods compute the similarity between two data points x and x'. Quantum kernel methods do the same, but the feature map $\phi(x)$ maps the data into a quantum Hilbert space.

$$K_Q(x, x') = |\langle \Phi(x) | \Phi(x') \rangle|^2$$

This is physically realized by:

1. Preparing a quantum state $|\Phi(x)\rangle$ using a parameterized circuit $U(x)$ applied to the vacuum state $|0\rangle^{\otimes n}$.
2. Applying the inverse circuit $U^\dagger(x')$ of the second data point.
3. Measuring the probability of finding the system in the all-zero state.

Certain quantum circuits, particularly those involving entanglement, produce feature maps that are hard to simulate classically. If the "decision boundary" between an FRB and RFI is simple in this high-dimensional quantum space (e.g., a linear hyperplane) but highly convoluted in classical space, the Quantum SVM will outperform the Classical SVM.[9] Specifically, "ZZ-Feature Maps" and "IQP Encoding" use two-qubit entangling gates to encode correlations between different features (e.g., the correlation between the 1st and 2nd order scattering coefficients). This allows the model to learn "texture" and "shape" dependencies that classical linear models miss.[25]

5.2 Architecture: The ZZ-Feature Map

For time-series anomaly detection, capturing the correlation between time-steps (or scale-steps in WST) is key. The **ZZ-Feature Map** is ideal for this.

The unitary operation is defined as:

$$U_\Phi(\mathbf{x}) = \exp\left(i\sum_j x_j Z_j + i\sum_{j,k}(\pi - x_j)(\pi - x_k) Z_j Z_k\right)$$

- The $x_j Z_j$ term encodes individual feature values (single-qubit rotation).
- The $Z_j Z_k$ term encodes the non-linear relationship (entanglement) between feature j and feature k.[27]

In the context of QETD, if x_j represents the energy at scale λ_1 and x_k represents the energy at scale λ_2, the quantum circuit naturally encodes the *inter-scale dependency*, precisely the information that defines a "chirp" or a "burst."

5.3 Quantum Anomaly Detection (Unsupervised)

Since we are looking for rare or unknown events, supervised classification (Signal vs. Noise) is risky, it requires training examples of signals we haven't found yet.

QETD will utilize **Quantum One-Class SVM (OC-SVM)** or **Quantum Autoencoders (QAE)**.

- **One-Class SVM:** The quantum model is trained *only* on the background noise (which is abundant). It learns to map the noise to a specific region of the Hilbert space. Any new signal that maps *outside* this region is flagged as an anomaly.[13]
- **Quantum Autoencoder:** The circuit tries to compress the input state into a smaller number of qubits and then reconstruct it. The "reconstruction fidelity" is the metric. Signals that are statistically different from the training noise will have poor reconstruction fidelity, triggering an alert.[3]

5.4 Hybrid Quantum-Classical Neural Networks (HQNN)

Given the current limitations of NISQ (Noisy Intermediate-Scale Quantum) hardware (limited qubits, noise), a fully quantum pipeline is not yet feasible for Tb/s data.

QETD proposes a **Hybrid Architecture**:

1. **Classical Head (CNN):** A lightweight classical network (CenterNet-based) performs initial reduction and candidate generation.
2. **Quantum Tail (VQC):** The difficult "edge cases", signals that are ambiguous to the classical net, are forwarded to the Quantum Processor (or simulator). The VQC acts as the ultimate arbiter, using its superior feature resolution to make the final decision.[29]

6. MLOps: Scale, Reliability, and Infrastructure

Deploying QETD is not just a physics problem; it is a systems engineering challenge of the highest order. It requires an operational framework capable of handling the SKA's data lifecycle.

6.1 The SKA Data Lifecycle

The SKA is a distributed instrument. Data flows from the dishes (Edge) to the Central Signal Processor (CSP), then to the Science Data Processor (SDP), and finally to the global network of SKA Regional Centres (SRCs).[5]

Stage	Data Rate / Vol	Compute Environment	QETD Role
Antenna/Station	~2 Tb/s	FPGA / ASIC	None (Data capture)
CSP (Correlator)	~8.9 Tb/s	FPGA / GPU	Coarse RFI flagging
SDP (Science Processor)	~1 TB/s	HPC Cluster	**Wavelet Preprocessing & Candidate Generation**
Regional Centre (SRC)	~300 PB/yr	Cloud / Grid	**Quantum Model Training & Archival Search**

6.2 MLOps Framework Components

To manage the lifecycle of the hybrid models, QETD will employ a strict MLOps stack compliant with SKA software standards.[32]

- **Data Version Control (DVC):** In science, reproducibility is paramount. DVC will be used to version control the datasets (e.g., "Simulated FRB Set v4.2" vs "Real Noise

Set v1.0"). This allows researchers to reproduce the exact training conditions of a deployed model.[33]

- **Kubeflow & Containerization:** The pipeline involves heterogeneous compute: CPUs for I/O, GPUs for Wavelets, and QPUs (or simulators) for classification. Kubeflow will orchestrate these as containerized microservices on Kubernetes. This ensures the pipeline is portable.
- **Model Registry (MLflow):** As we experiment with different quantum circuits (Z-map vs ZZ-map, varying entanglement depth), MLflow will log the hyperparameters and performance metrics (Recall, Precision). This allows for systematic benchmarking of "Quantum Advantage".[33]

6.3 The "Inference at the Edge" Strategy

Latency is the enemy. We cannot send terabits of data to a cloud-based quantum computer.

Strategy:

1. **Offline Training:** The heavy lifting (finding the optimal quantum kernel parameters) happens offline at the SRCs using archival data.
2. **Model Distillation / Quantization:** The trained quantum model is either distilled into a classical neural network (student-teacher model) or executed on a localized, dedicated quantum co-processor (if available in future hardware generations).
3. **Edge Deployment:** The lightweight inference engine runs on the SDP nodes, processing the Wavelet coefficients in real-time (< 10 seconds latency).[4]

7. Mathematical Formalism and Deep Technical Description

7.1 Wavelet Scattering Properties

Let $\psi(t)$ be a complex wavelet with center frequency λ. The transform is $Wx(t, \lambda) = x * \psi_\lambda(t)$.

The WST constructs a representation $\Phi(x)$ that is Lipschitz continuous to deformations.

If $x_\tau(t) = x(t - \tau(t))$ is a deformed signal (where $\tau(t)$ is small), then:|

$$\|\Phi(x) - \Phi(x_\tau)\| \leq C\|\nabla\tau\|_\infty\|x\|$$

This inequality guarantees that small changes in the signal shape (due to noise or slight DM variations) result in small changes in the feature vector, whereas the "distance" to RFI remains large.[12]

7.2 Quantum State Preparation: The Input Problem

A major bottleneck in QML is loading classical data x into the quantum state $|\psi\rangle$.

- **Amplitude Encoding:** Encodes N features into $\log_2 N$ qubits ($x = \sum x_i |i\rangle$). This is space-efficient but requires deep circuits ($O(N)$ depth).
- **Angle Encoding:** Encodes N features into N qubits ($|\psi\rangle = \bigotimes \cos(x_i)|0\rangle + \sin(x_i)|1\rangle$). This is depth-efficient ($O(1)$) but uses more qubits. **QETD Choice:** Given that the Wavelet Scattering Transform compresses the input to a small vector (approx 50-60 dimensions), we can utilize **Angle Encoding** or **Dense Angle Encoding** (2 features per qubit) which is feasible on current 50-100 qubit devices (IBM Eagle/Heron).[25]

7.3 The Quantum Kernel Matrix

The kernel matrix K used in the SVM is estimated by the "Swap Test" or "Inversion Test".

For the ZZ-feature map $\mathcal{U}_\Phi(\mathbf{x})$, the kernel entry is:

$$K(\mathbf{x}_i, \mathbf{x}_j) = |\langle 0|\mathcal{U}_\Phi^\dagger(\mathbf{x}_i)\mathcal{U}_\Phi(\mathbf{x}_j)|0\rangle|^2$$

The optimization problem (finding the support vectors) is convex:

$$\min_\alpha \frac{1}{2}\alpha^T Q \alpha - \mathbf{1}^T \alpha$$

where $Q_{ij} = y_i y_j K(\mathbf{x}_i, \mathbf{x}_j)$. Because this is convex, it guarantees finding the *global* optimum separation, unlike neural networks which can get stuck in local minima.[38]

8. Things to Explore and Future Horizons

The QETD framework serves as a foundation for broader inquiries into quantum-enhanced science.

8.1 Neuromorphic Quantum Computing

Given the strong parallels between biological neural spikes and astrophysical bursts, future research should explore **Spiking Quantum Neural Networks (SQNNs)**. These models encode information in the *timing* of quantum pulses rather than amplitudes, potentially offering orders-of-magnitude energy reduction, crucial for space-based telescopes where power is limited.[39]

8.2 The Search for Technosignatures (SETI)

The ultimate "anomaly" is a signal from an extraterrestrial intelligence. SETI signals might be encoded to look like noise to a simple observer (spread spectrum). Quantum kernels, capable of detecting high-order entropy anomalies and non-random entanglement structures, might reveal a "signal" where classical FFTs see only static. This represents a new frontier for SETI: Quantum SETI.[40]

8.3 Real-Time Feedback Loops

Currently, detection is passive. In the future, a fast QETD trigger could actively control the telescope, commanding it to freeze buffers, dump raw voltage data, or even slew the dishes to track a repeating source. This requires pushing quantum latency down to the millisecond regime, potentially requiring hybrid Quantum-FPGA interconnects.

9. Conclusion

The data intensity of the Exascale Era presents an existential challenge to astronomy. We are building eyes capable of seeing the edge of the universe, but our brains (classical computers) are too slow to process what we see. **Quantum-Enhanced Transient Discovery** offers a path forward.

By synergizing the mathematical rigour of Wavelet Scattering with the high-dimensional expressivity of Quantum Machine Learning, QETD creates a detection filter that is both robust to the chaotic noise of the Anthropocene and sensitive to the faint whispers of the cosmos. It moves the problem from the brute-force template matching of the 20th century to the intelligent, probabilistic anomaly detection of the 21st.

While significant engineering hurdles remain, particularly regarding NISQ device noise and data loading speeds, the potential rewards are transformative. A successful QETD system would not only increase the FRB detection rate by orders of magnitude but could also unveil entirely new classes of astrophysical phenomena that have, until now, remained hidden in the noise.

Works cited

1. Use case class 5: real-time challenges, data irreversibility - PUNCH4NFDI, accessed on February 16, 2026, https://www.punch4nfdi.de/use_cases/use_case_class_5/
2. DRAFTS: A Deep Learning-Based Radio Fast Transient Search Pipeline - arXiv, accessed on February 16, 2026, https://arxiv.org/html/2410.03200v1
3. Towards an anomaly detection pipeline for gravitational ... - arXiv.org, accessed on February 16, 2026, https://arxiv.org/abs/2511.13154
4. Square Kilometre Array Computational Challenges - CERN Indico, accessed on February 16, 2026,

https://indico.cern.ch/event/184526/contributions/322917/attachments/253364/354189/PAlexander_SKA_Computational_Challenges.pdf
5. Big telescope, big data: towards exascale with the Square Kilometre ..., accessed on February 16, 2026, https://pmc.ncbi.nlm.nih.gov/articles/PMC7015288/
6. The bursting sky | SKAO, accessed on February 16, 2026, https://www.skao.int/en/explore/science-goals/146/bursting-sky
7. Radio Astronomy Research Articles - Page 5 | R Discovery, accessed on February 16, 2026, https://discovery.researcher.life/topic/radio-astronomy/27303557?page=5
8. DRAFTS: A Deep-learning-based Radio Fast Transient Search Pipeline - ResearchGate, accessed on February 16, 2026, https://www.researchgate.net/publication/387753561_DRAFTS_A_Deep-learning-based_Radio_Fast_Transient_Search_Pipeline
9. Periodic Unitary Encoding for Quantum Anomaly Detection of Temporal Series - SciTePress, accessed on February 16, 2026, https://www.scitepress.org/Papers/2025/135378/135378.pdf
10. Hybrid Quantum Machine Learning for Anomaly Detection in Critical, accessed on February 16, 2026, https://gjeta.com/sites/default/files/GJETA-2025-0198.pdf
11. Analyzing the 21cm forest with Wavelet Scattering Transform: Insight into non-Gaussian features of the 21cm forest - arXiv, accessed on February 16, 2026, https://arxiv.org/html/2504.14656v2
12. Wavelet Scattering - MATLAB & Simulink - MathWorks, accessed on February 16, 2026, https://www.mathworks.com/help/wavelet/ug/wavelet-scattering.html
13. Quantum Kernel Anomaly Detection, Using AR Features, Achieves Non-Contact Monitoring For Smart Manufacturing Equipment, accessed on February 16, 2026, https://quantumzeitgeist.com/quantum-detection-features-kernel-anomaly-achieves-non-contact-monitoring-smart-manufacturing/
14. Quantum deep learning-based anomaly detection for enhanced network security, accessed on February 16, 2026, https://www.researchgate.net/publication/380293282_Quantum_deep_learning-based_anomaly_detection_for_enhanced_network_security
15. Radio Astronomy Fundamentals II John Simonetti Spring 2010 The “signal” from an astronomical radio source is hard to distin, accessed on February 16, 2026, https://www1.phys.vt.edu/~jhs/phys3154/Radio%20Astronomy%20Fun%202.pdf
16. A Self-learning Neural Network Approach for Radio Frequency Interference Detection and Removal in Radio Astronomy - OSTI, accessed on February 16, 2026, https://www.osti.gov/servlets/purl/1906939
17. Noise characteristics in spaceflight multichannel EEG | PLOS One - Research journals, accessed on February 16, 2026, https://journals.plos.org/plosone/article?id=10.1371/journal.pone.0280822
18. Comparison of STFT and wavelet transform methods in determining epileptic seizure activity in EEG signals for real-time application - PubMed, accessed on February 16, 2026, https://pubmed.ncbi.nlm.nih.gov/15809098/

19. Comparison of Wavelet Transform and FFT Methods in the Analysis of EEG Signals - PubMed, accessed on February 16, 2026, https://pubmed.ncbi.nlm.nih.gov/12018610/
20. Time-Frequency Analysis and Wavelets - Sapien Labs | Shaping the Future of Mind Health, accessed on February 16, 2026, https://sapienlabs.org/time-frequency-analysis-and-wavelets/
21. Comparing EEG-Based Epilepsy Diagnosis Using Neural Networks and Wavelet Transform, accessed on February 16, 2026, https://www.mdpi.com/2076-3417/13/18/10412
22. Scattering transforms, recent applications in astrophysics and cosmology - SISSA Indico server, accessed on February 16, 2026, https://indico.sissa.it/event/77/contributions/1512/attachments/352/445/EA_IFPU_29_06.pdf
23. A New Wavelet Scattering Transform-Based Statistic for Cosmological Analysis of Large-Scale Structure - arXiv, accessed on February 16, 2026, https://arxiv.org/html/2505.14400v1
24. Analyzing the 21-cm forest with wavelet scattering transform: Insight into non-Gaussian features of the 21-cm forest | Request PDF - ResearchGate, accessed on February 16, 2026, https://www.researchgate.net/publication/395962506_Analyzing_the_21-cm_forest_with_wavelet_scattering_transform_Insight_into_non-Gaussian_features_of_the_21-cm_forest
25. Shallow entangled circuits for quantum time series prediction on IBM devices - PMC, accessed on February 16, 2026, https://pmc.ncbi.nlm.nih.gov/articles/PMC12700859/
26. Quantum-UnIMP: Hybrid Quantum-Classical Imputation - Emergent Mind, accessed on February 16, 2026, https://www.emergentmind.com/topics/quantum-unimp
27. Hardware-Agnostic Quantum Kernel Feature Mapping for Anomaly Detection in Critical Infrastructure: A Cross-Testbed Validation on NISQ Processors | Sciety, accessed on February 16, 2026, https://sciety.org/articles/activity/10.21203/rs.3.rs-8682004/v1
28. Implementation and Performance Evaluation of Quantum Machine Learning Algorithms for Binary Classification - Bournemouth University, accessed on February 16, 2026, https://eprints.bournemouth.ac.uk/40696/1/software-03-00024.pdf
29. An In-Depth Comparative Study of Quantum-Classical Encoding Methods for Network Intrusion Detection - IEEE Xplore, accessed on February 16, 2026, https://ieeexplore.ieee.org/iel8/8782661/8901158/10869381.pdf
30. Quantum-Enhanced Representation Learning: A Quanvolutional Autoencoder Approach against DDoS Threats - MDPI, accessed on February 16, 2026, https://www.mdpi.com/2504-4990/6/2/44
31. SKAO / ska-data-lifecycle - GitLab, accessed on February 16, 2026, https://gitlab.com/ska-telescope/ska-data-lifecycle
32. SKA Data Lifecycle Management, accessed on February 16, 2026, https://developer.skao.int/projects/ska-data-lifecycle/en/latest/

33. From Data to Deployment: Implementing MLOps with Airflow, DVC, Mlflow and Kubernetes | by Syeda Areeba Nadeem | Medium, accessed on February 16, 2026, https://medium.com/@syeda.areeba.nadeem/from-data-to-deployment-implementing-mlops-with-airflow-dvc-mlflow-and-kubernetes-6b2df084afba
34. MLOps Archives - DVC, accessed on February 16, 2026, https://dvc.org/blog/tags/mlops/
35. Machine Learning Operations | MLOps Tutorial | DVC | MLFlow | KubeFlow - YouTube, accessed on February 16, 2026, https://www.youtube.com/watch?v=EQ0zc2VRX1w
36. Deep Scattering Spectrum, accessed on February 16, 2026, https://www.di.ens.fr/~mallat/papiers/AudioScatSpectrum.pdf
37. [2102.07533] Low-depth Quantum State Preparation - arXiv.org, accessed on February 16, 2026, https://arxiv.org/abs/2102.07533
38. Practical Evaluation of Quantum Kernel Methods for Radar Micro-Doppler Classification on Noisy Intermediate-Scale Quantum (NISQ) Hardware - arXiv, accessed on February 16, 2026, https://arxiv.org/html/2601.22194v1
39. RFI Detection with Spiking Neural Networks - arXiv, accessed on February 16, 2026, https://arxiv.org/html/2311.14303v2
40. Scientific Publications - FAST - Five-hundred-meter Aperture Spherical radio Telescope, accessed on February 16, 2026, https://fast.bao.ac.cn/cms/article/91/